

# Twin-Delayed DDPG (TD3) with Bipedal Walker

Chien-Te Lee

2020/8/20

# Outline

---

- **OpenAI Gym environment**
- **Model Architecture**
- **Algorithm**
- **Result**
- **Reference**

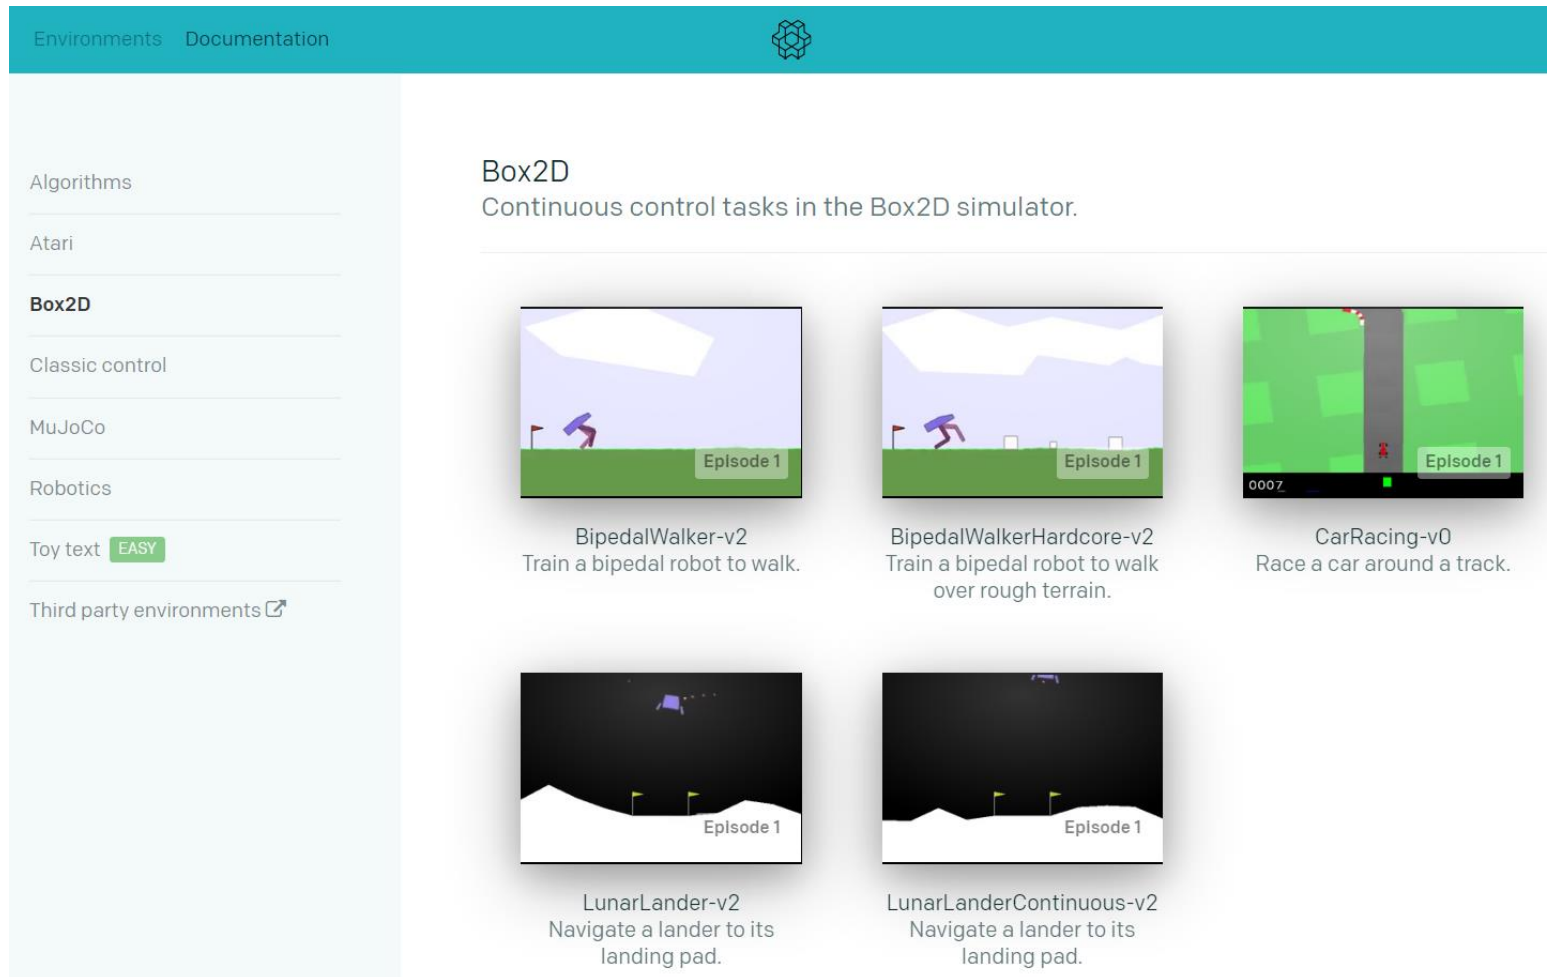
# Outline

---

- **OpenAI Gym environment**
- **Model Architecture**
- **Algorithm**
- **Result**
- **Reference**

# OpenAI Gym environment


- OpenAI Gym is a toolkit for developing and comparing reinforcement learning (RL) algorithms.



# OpenAI Gym environment



---

- We are going to use BipedalWalker-v3 environment.
- The environment simulates a random and uneven terrain.
- The goal is to develop an RL algorithm which can drive the Bipedal Walker, to walk to the end of the terrain without falling down.



# OpenAI Gym environment

- The 24 states represent the x and y velocities, the joint angle and angular velocities, ground contact of legs, and lidar readings for ground distance.
- The 4 actions represent the torque control of 2 hips and 2 knees.



# Outline

---


- **OpenAI Gym environment**
- **Model Architecture**
- **Algorithm**
- **Result**
- **Reference**

# Model Architecture

---

- Twin-Delayed DDPG contains 2 actor (policy) networks: actor-evaluate and actor-target. They have same architecture and different random weights.
- The actor (policy) takes the states and predicts what actions to perform.
- Notation:  $a = \pi(s)$

## Actor Network Architecture






# Model Architecture

- Twin-Delayed DDPG contains 4 critic networks: 2 critic-evaluate and 2 critic-target. They have same architecture and different random weights.
- Two critics take the states and actions to predict estimated q-value.
- Notation: q-value =  $Q(s, a)$

## Critic Network Architecture



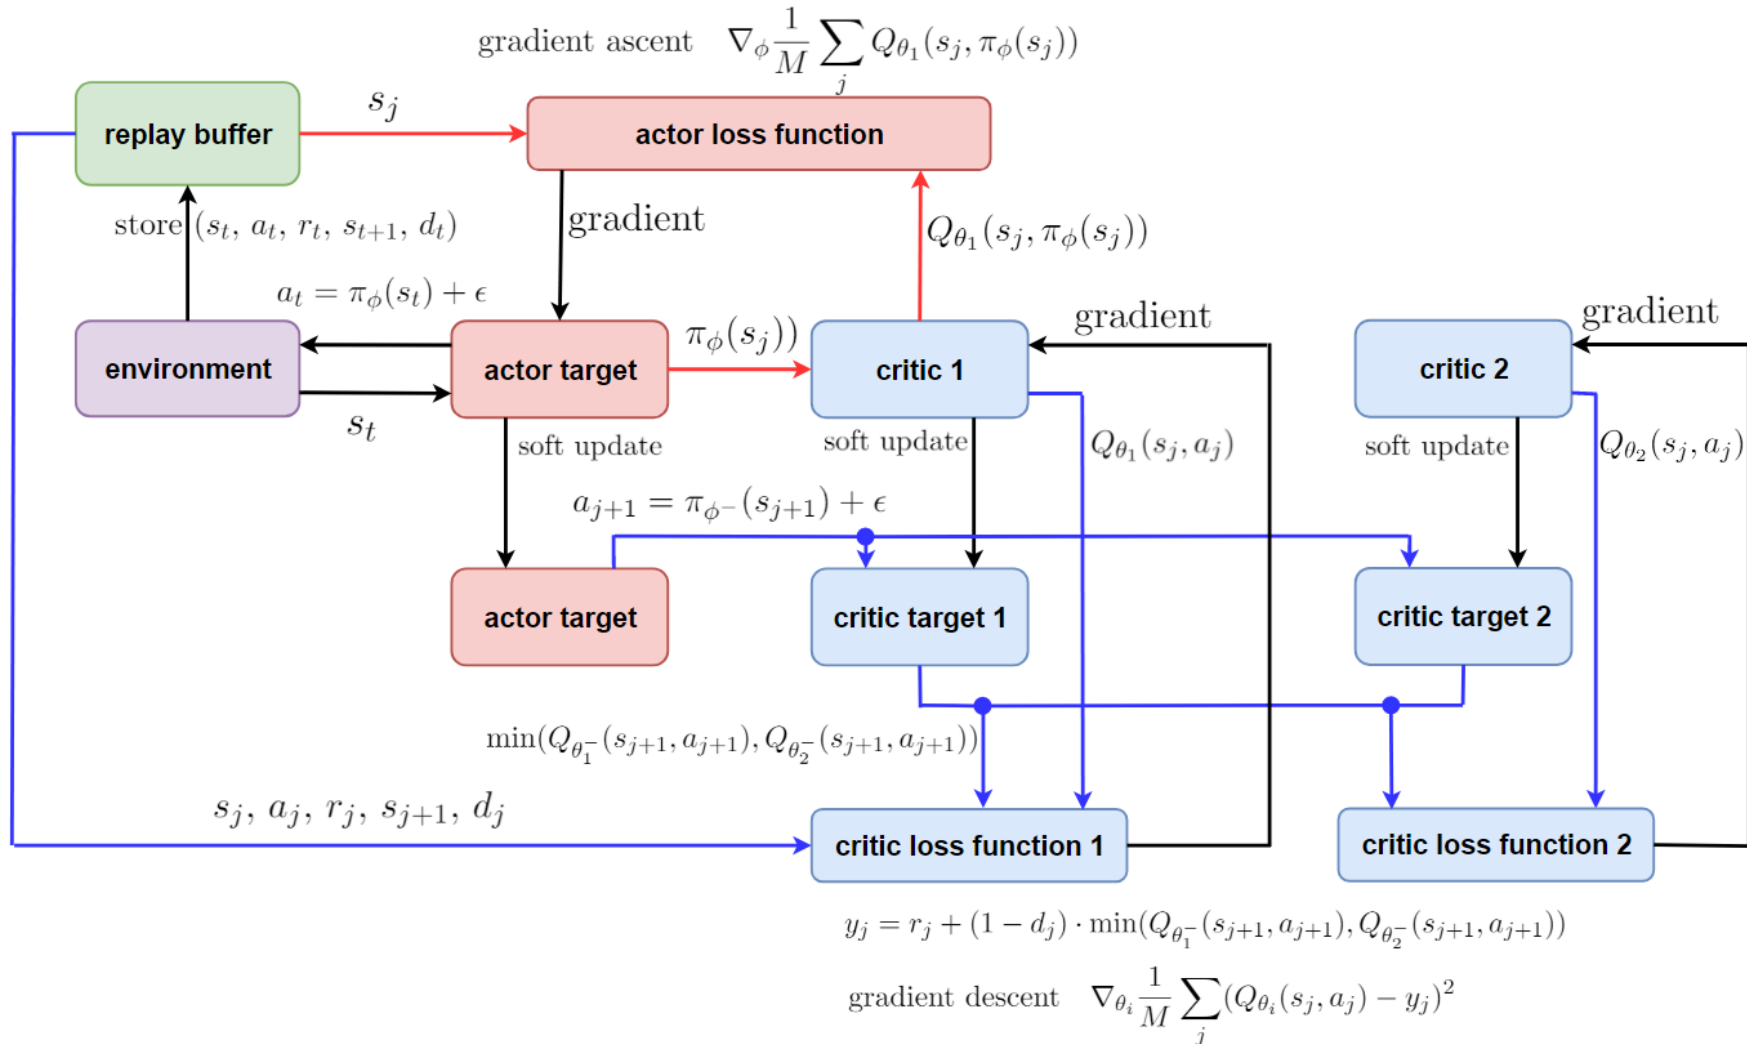
# Outline

---

- **OpenAI Gym environment**
- **Model Architecture**
- **Algorithm Implementation**
- **Result**
- **Reference**

# Algorithm

**Twin-Delayed DDPG Flow Diagram**



# Algorithm

---

---

**Algorithm 1** Twin Delayed DDPG

---

```
1: Initialize two critic nets  $Q$  with random weights  $\theta_1, \theta_2$ , one actor net  $\pi$  with random weights  $\phi$ 
2: Set target net weights  $\theta_1^- \leftarrow \theta_1, \theta_2^- \leftarrow \theta_2, \phi^- \leftarrow \phi$ 
3: Initialize learning rate  $\eta$ , reward decay  $\gamma$ , soft update parameter  $\tau$ 
4: Initialize noise standard deviation  $\sigma$ , noise boundary  $c$ , delay interval  $K$ 
5: Initialize replay memory  $R$  to capacity  $N$  and minibatch size  $M$ 
6: Initialize episode  $E$  with play time  $T$ 
7: for  $episode = 1$  to  $E$  do
8:   Reset environment and agent to random initial state  $s_0$ 
9:   Set done  $d = 0$ 
10:  for  $t = 1$  to  $T$  do
11:    Select action with exploration noise  $a_t = \pi_\phi(s_t) + \epsilon$ , and  $\epsilon \sim clip(\mathcal{N}(0, \sigma), -c, c)$ 
12:    Executes action  $a_t$  at state  $s_t$  and observes reward  $r_t$ , next state  $s_{t+1}$ , and done  $d_t$ 
13:    Store transition  $(s_t, a_t, r_t, s_{t+1}, d_t)$  in replay memory  $R$ 
14:    Sample minibatch with  $M$  transitions  $(s_j, a_j, r_j, s_{j+1}, d_j)$  from replay memory  $R$ 
15:    Calculate next actions for minibatch  $a_{j+1} = \pi_{\phi^-}(s_{j+1}) + \epsilon$ , and  $\epsilon \sim clip(\mathcal{N}(0, \sigma), -c, c)$ 
16:    Calculate target  $y_j = r_j + (1 - d_j) \cdot \gamma \cdot \min(Q_{\theta_1^-}(s_{j+1}, a_{j+1}), Q_{\theta_2^-}(s_{j+1}, a_{j+1}))$ 
17:    Perform two gradient descents to update two critic nets respectively:
```

$$\nabla_{\theta_i} \frac{1}{M} \sum_j (Q_{\theta_i}(s_j, a_j) - y_j)^2 \quad \text{for } i \in \{1, 2\}, j \text{ in minibatch}$$

```
18:    if  $t \bmod K = 0$  then
19:      Perform gradient ascent to update actor net using critic net 1:
```

$$\nabla_{\phi} \frac{1}{M} \sum_j Q_{\theta_1}(s_j, \pi_{\phi}(s_j)) \quad \text{for } j \text{ in minibatch}$$

```
20:      Soft update target nets:
```

$$\begin{aligned} \theta_i^- &\leftarrow \tau \cdot \theta_i + (1 - \tau) \cdot \theta_i^- & \text{for } i \in \{1, 2\} \\ \phi^- &\leftarrow \tau \cdot \phi + (1 - \tau) \cdot \phi^- \end{aligned} .$$

```
21:    end if
22:    if done  $d = 1$  then
23:      Break
24:    end if
25:    Move on to next state  $s_t \leftarrow s_{t+1}$ 
26:  end for
27: end for
```

---

# Algorithm

---

## Algorithm 1 Twin Delayed DDPG

- 1: Initialize two critic nets  $Q$  with random weights  $\theta_1, \theta_2$ , one actor net  $\pi$  with random weights  $\phi$
- 2: Set target net weights  $\theta_1^- \leftarrow \theta_1, \theta_2^- \leftarrow \theta_2, \phi^- \leftarrow \phi$
- 3: Initialize learning rate  $\eta$ , reward decay  $\gamma$ , soft update parameter  $\tau$
- 4: Initialize noise standard deviation  $\sigma$ , noise boundary  $c$ , delay interval  $K$
- 5: Initialize replay memory  $R$  to capacity  $N$  and minibatch size  $M$
- 6: Initialize episode  $E$  with play time  $T$
- 7: **for**  $episode = 1$  to  $E$  **do**
- 8:     Reset environment and agent to random initial state  $s_0$
- 9:     Set done  $d = 0$
- 10:    **for**  $t = 1$  to  $T$  **do**
- 11:       Select action with exploration noise  $a_t = \pi_\phi(s_t) + \epsilon$ , and  $\epsilon \sim clip(\mathcal{N}(0, \sigma), -c, c)$
- 12:       Executes action  $a_t$  at state  $s_t$  and observes reward  $r_t$ , next state  $s_{t+1}$ , and done  $d_t$
- 13:       Store transition  $(s_t, a_t, r_t, s_{t+1}, d_t)$  in replay memory  $R$

1. Initialize random weights for 2 actor networks and 4 critic networks.

2. Set parameters such as learning rate, reward decay, or soft update parameter.

# Algorithm

---

## Algorithm 1 Twin Delayed DDPG

---

- 1: Initialize two critic nets  $Q$  with random weights  $\theta_1, \theta_2$ , one actor net  $\pi$  with random weights  $\phi$
- 2: Set target net weights  $\theta_1^- \leftarrow \theta_1, \theta_2^- \leftarrow \theta_2, \phi^- \leftarrow \phi$
- 3: Initialize learning rate  $\eta$ , reward decay  $\gamma$ , soft update parameter  $\tau$
- 4: Initialize noise standard deviation  $\sigma$ , noise boundary  $c$ , delay interval  $K$
- 5: Initialize replay memory  $R$  to capacity  $N$  and minibatch size  $M$
- 6: Initialize episode  $E$  with play time  $T$
- 7: **for**  $episode = 1$  to  $E$  **do**
- 8:     Reset environment and agent to random initial state  $s_0$
- 9:     Set done  $d = 0$
- 10:    **for**  $t = 1$  to  $T$  **do**
- 11:       Select action with exploration noise  $a_t = \pi_\phi(s_t) + \epsilon$ , and  $\epsilon \sim clip(\mathcal{N}(0, \sigma), -c, c)$
- 12:       Executes action  $a_t$  at state  $s_t$  and observes reward  $r_t$ , next state  $s_{t+1}$ , and done  $d_t$
- 13:       Store transition  $(s_t, a_t, r_t, s_{t+1}, d_t)$  in replay memory  $R$

1. Actor-evaluate selects action with bounded random noise.
2. Execute action at state in gym.
3. Observe reward and next state.
4. Store transition in replay memory.

# Algorithm

---

- 14: Sample minibatch with  $M$  transitions  $(s_j, a_j, r_j, s_{j+1}, d_j)$  from replay memory  $R$
- 15: Calculate next actions for minibatch  $a_{j+1} = \pi_{\phi^-}(s_{j+1}) + \epsilon$ , and  $\epsilon \sim \text{clip}(\mathcal{N}(0, \sigma), -c, c)$
- 16: Calculate target  $y_j = r_j + (1 - d_j) \cdot \gamma \cdot \min(Q_{\theta_1^-}(s_{j+1}, a_{j+1}), Q_{\theta_2^-}(s_{j+1}, a_{j+1}))$
- 17: ~~Perform two gradient descents to update two critic nets respectively.~~

$$\nabla_{\theta_i} \frac{1}{M} \sum_j (Q_{\theta_i}(s_j, a_j) - y_j)^2 \quad \text{for } i \in \{1, 2\}, j \text{ in minibatch}$$

1. Sample minibatch transitions for experience replay.
2. Use actor-target to calculate next actions with bounded random noise for minibatch.
3. Calculate 2 next q-values with 2 critic-targets. Select the smaller next q-value of the two, and combine with the reward to calculate target for each transition in the minibatch.

# Algorithm

- 14: Sample minibatch with  $M$  transitions  $(s_j, a_j, r_j, s_{j+1}, d_j)$  from replay memory  $R$   
15: Calculate next actions for minibatch  $a_{j+1} = \pi_{\phi^-}(s_{j+1}) + \epsilon$ , and  $\epsilon \sim \text{clip}(\mathcal{N}(0, \sigma), -c, c)$   
16: Calculate target  $y_j = r_j + (1 - d_j) \cdot \gamma \cdot \min(Q_{\theta_1^-}(s_{j+1}, a_{j+1}), Q_{\theta_2^-}(s_{j+1}, a_{j+1}))$   
17: Perform two gradient descents to update two critic nets respectively:

$$\nabla_{\theta_i} \frac{1}{M} \sum_j (Q_{\theta_i}(s_j, a_j) - y_j)^2 \quad \text{for } i \in \{1, 2\}, j \text{ in minibatch}$$

- 18: **if**  $t \bmod K = 0$  **then**  
19:     Perform gradient ascent to update actor net using critic net 1:

$$\nabla_{\phi} \frac{1}{M} \sum_j Q_{\theta_1}(s_j, \pi_{\phi}(s_j)) \quad \text{for } j \text{ in minibatch}$$

1. For every experience replay, perform gradient descend and update two critic-evaluate with the targets calculated in previous slide.
2. For every  $K$  experience replay, perform gradient ascend and update actor-evaluate using the “first” critic-evaluate.
3. The actor update is “delayed” for stabilizing training process.



# Algorithm

---

18:       **if**  $t \bmod K = 0$  **then**  
19:           Perform gradient ascent to update actor net using critic net 1:

$$\nabla_{\phi} \frac{1}{M} \sum_j Q_{\theta_1}(s_j, \pi_{\phi}(s_j)) \quad \text{for } j \text{ in minibatch}$$

20:           Soft update target nets:

$$\begin{aligned} \theta_i^- &\leftarrow \tau \cdot \theta_i + (1 - \tau) \cdot \theta_i^- & \text{for } i \in \{1, 2\} \\ \phi^- &\leftarrow \tau \cdot \phi + (1 - \tau) \cdot \phi^- \end{aligned}$$

21:       **end if**  
22:       **if** done  $d = 1$  **then**  
23:           Break  
24:       **end if**  
25:       Move on to next state  $s_t \leftarrow s_{t+1}$   
26:   **end for**  
27: **end for**

---

1. Every K experience replay, when we update actor-evaluate, we also do soft update for actor and critic target networks.

2. This time-step finishes, this state move on to next state.


# Outline

---

- **OpenAI Gym environment**
- **Model Architecture**
- **Algorithm Implementation**
- **Result**
- **Reference**

# Result

- The training process is shown below. After training, the average score for greedy agent (action without noise) is around 328, which is considered solved for this environment.



# Result

---

- We can see that the Bipedal Walker can successfully walk to the end of the terrain.



# Outline

---

- **OpenAI Gym environment**
- **Model Architecture**
- **Algorithm Implementation**
- **Result**
- **Reference**

# Reference

---

- [Continuous control with deep reinforcement learning](#)
- [Addressing Function Approximation Error in Actor-Critic Methods](#)
- [Deep Deterministic Policy Gradient](#)
- [Deep Reinforcement Learning. Deep Deterministic Policy Gradient \(DDPG\) algorithm](#)
- [Deep Deterministic Policy Gradient \(DDPG\)](#)
- [Twin Delayed DDPG](#)
- [TD3: Learning To Run With AI](#)

Thank you for your attention!!